

CHURN PREDICTION

Importing Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
%matplotlib inline
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Dropout
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import warnings
warnings.filterwarnings('ignore')
```

Loading Datasets

```
In [2]: client_data = pd.read_csv("../Data/client_data.csv")
client_data.head()
```

```
Out[2]:
```

	id	channel_sales	cons_12m	cons_gas_12m	cons_last_month	date_activ
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsosbicdxkicaua	0	54946	0	2013-06-15
1	d29c2c54acc38ff3c0614d0a653813dd	MISSING	4660	0	0	2009-08-21
2	764c75f661154dac3a6c254cd082ea7d	foosdfpfkusacimwkcsosbicdxkicaua	544	0	0	2010-04-16
3	bba03439a292a1e166f80264c16191cb	lmkebamcaaclubfxadlmueccxoimlema	1584	0	0	2010-03-30
4	149d57cf92fc41cf94415803a877cb4b	MISSING	4425	0	526	2010-01-13

5 rows × 26 columns

```
In [3]: price_data = pd.read_csv("../Data/price_data.csv")
price_data.head()
```

```
Out[3]:
```

	id	price_date	price_off_peak_var	price_peak_var	price_mid_peak_var	price_off_peak_fix	price
0	038af19179925da21a25619c5a24b745	2015-01-01	0.151367	0.0	0.0	44.266931	
1	038af19179925da21a25619c5a24b745	2015-02-01	0.151367	0.0	0.0	44.266931	
2	038af19179925da21a25619c5a24b745	2015-03-01	0.151367	0.0	0.0	44.266931	
3	038af19179925da21a25619c5a24b745	2015-04-01	0.149626	0.0	0.0	44.266931	
4	038af19179925da21a25619c5a24b745	2015-05-01	0.149626	0.0	0.0	44.266931	

Data Description

```
In [4]: client_data.shape
```

```
Out[4]: (14606, 26)
```

```
In [5]: client_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14606 entries, 0 to 14605
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    14606 non-null  object
1   channel_sales                        14606 non-null  object
2   cons_12m                             14606 non-null  int64
3   cons_gas_12m                         14606 non-null  int64
4   cons_last_month                     14606 non-null  int64
5   date_activ                           14606 non-null  object
6   date_end                             14606 non-null  object
7   date_modif_prod                     14606 non-null  object
8   date_renewal                         14606 non-null  object
9   forecast_cons_12m                   14606 non-null  float64
10  forecast_cons_year                   14606 non-null  int64
11  forecast_discount_energy             14606 non-null  float64
12  forecast_meter_rent_12m              14606 non-null  float64
13  forecast_price_energy_off_peak       14606 non-null  float64
14  forecast_price_energy_peak           14606 non-null  float64
15  forecast_price_pow_off_peak          14606 non-null  float64
16  has_gas                              14606 non-null  object
17  imp_cons                             14606 non-null  float64
18  margin_gross_pow_ele                 14606 non-null  float64
19  margin_net_pow_ele                   14606 non-null  float64
20  nb_prod_act                          14606 non-null  int64
21  net_margin                           14606 non-null  float64
22  num_years_antig                      14606 non-null  int64
23  origin_up                            14606 non-null  object
24  pow_max                              14606 non-null  float64
25  churn                                14606 non-null  int64
dtypes: float64(11), int64(7), object(8)
memory usage: 2.9+ MB
```

```
In [6]: price_data.shape
```

Out[6]: (193002, 8)

```
In [7]: price_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193002 entries, 0 to 193001
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    193002 non-null  object
1   price_date                           193002 non-null  object
2   price_off_peak_var                   193002 non-null  float64
3   price_peak_var                       193002 non-null  float64
4   price_mid_peak_var                   193002 non-null  float64
5   price_off_peak_fix                   193002 non-null  float64
6   price_peak_fix                       193002 non-null  float64
7   price_mid_peak_fix                   193002 non-null  float64
dtypes: float64(6), object(2)
memory usage: 11.8+ MB
```

Data Statistics

```
In [8]: client_data.describe()
```

	cons_12m	cons_gas_12m	cons_last_month	forecast_cons_12m	forecast_cons_year	forecast_discount_energy	forecast_
count	1.460600e+04	1.460600e+04	14606.000000	14606.000000	14606.000000	14606.000000	
mean	1.592203e+05	2.809238e+04	16090.269752	1868.614880	1399.762906	0.966726	
std	5.734653e+05	1.629731e+05	64364.196422	2387.571531	3247.786255	5.108289	
min	0.000000e+00	0.000000e+00	0.000000	0.000000	0.000000	0.000000	
25%	5.674750e+03	0.000000e+00	0.000000	494.995000	0.000000	0.000000	
50%	1.411550e+04	0.000000e+00	792.500000	1112.875000	314.000000	0.000000	
75%	4.076375e+04	0.000000e+00	3383.000000	2401.790000	1745.750000	0.000000	
max	6.207104e+06	4.154590e+06	771203.000000	82902.830000	175375.000000	30.000000	

```
In [9]: price_data.describe()
```

Out[9]:	price_off_peak_var	price_peak_var	price_mid_peak_var	price_off_peak_fix	price_peak_fix	price_mid_peak_fix
count	193002.000000	193002.000000	193002.000000	193002.000000	193002.000000	193002.000000
mean	0.141027	0.054630	0.030496	43.334477	10.622875	6.409984
std	0.025032	0.049924	0.036298	5.410297	12.841895	7.773592
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.125976	0.000000	0.000000	40.728885	0.000000	0.000000
50%	0.146033	0.085483	0.000000	44.266930	0.000000	0.000000
75%	0.151635	0.101673	0.072558	44.444710	24.339581	16.226389
max	0.280700	0.229788	0.114102	59.444710	36.490692	17.458221

Checking Null Values

```
In [10]: client_data.isna().sum()
```

```
Out[10]: id                                0
channel_sales                             0
cons_12m                                  0
cons_gas_12m                              0
cons_last_month                           0
date_activ                                0
date_end                                  0
date_modif_prod                           0
date_renewal                              0
forecast_cons_12m                         0
forecast_cons_year                        0
forecast_discount_energy                  0
forecast_meter_rent_12m                   0
forecast_price_energy_off_peak             0
forecast_price_energy_peak                 0
forecast_price_pow_off_peak                0
has_gas                                    0
imp_cons                                  0
margin_gross_pow_ele                       0
margin_net_pow_ele                         0
nb_prod_act                                0
net_margin                                 0
num_years_antig                            0
origin_up                                  0
pow_max                                    0
churn                                       0
dtype: int64
```

```
In [11]: price_data.isna().sum()
```

```
Out[11]: id                                0
price_date                                0
price_off_peak_var                        0
price_peak_var                            0
price_mid_peak_var                        0
price_off_peak_fix                        0
price_peak_fix                            0
price_mid_peak_fix                        0
dtype: int64
```


EDA

Conducted a thorough investigation of the dataset using Tableau. The key visualizations and insights from this Tableau-based EDA are embedded below to provide context and support our subsequent modeling efforts.

Churn Percent

```
In [12]: img = mpimg.imread("./images/Churn Percent.png")
plt.figure(figsize=(18,8))
plt.imshow(img)
plt.axis('off')
plt.show()
```

Churn Percent




From the above chart we can see that approximately 10% of the customers are switching

Channel Wise Churn Percent

```
In [13]: img = mpimg.imread("./images/Channel_sales_churn.png")
plt.figure(figsize=(18,8))
plt.imshow(img)
plt.axis('off')
plt.show()
```


Channel_sales_churn



From the above chart we can see that maximum customers from the "foosdfpfkusacimwkcsosbicdxkicaa" and "usilxuppasemublopkaaafesmlibmsdf" channel

Consumption based Churn Percent

```
In [14]: img = mpimg.imread("./images/Consumption.png")
plt.figure(figsize=(18,12))
plt.imshow(img)
plt.axis('off')
plt.show()
```




From the above chart we can see that customers having consumption between 0 and 100k are switching.

Gas Consumer Churn Percent

```
In [15]: img = mpimg.imread("./images/Gas_Consumer.png")
plt.figure(figsize=(10,8))
plt.imshow(img)
plt.axis('off')
plt.show()
```

```
plt.show()
```


Gas_Consumer



From this chart we can see that customers who take electricity but does not take gas supply has the higher percent of switching.

Subscribed Power Churn Percent

```
In [16]: img = mpimg.imread("./images/Power_subscribed.png")
plt.figure(figsize=(18,8))
plt.imshow(img)
plt.axis('off')
plt.show()
```



From this chart we can see that customers with subscribed power between 10 to 20 are switching.

Feature Engineering

```
In [17]: # Converting date columns to datetime format
date_cols = ['date_activ', 'date_end', 'date_renewal', 'date_modif_prod']
for col in date_cols:
```

```
client_data[col] = pd.to_datetime(client_data[col])

price_data["price_date"] = pd.to_datetime(price_data["price_date"])
```

Difference between off-peak prices in December and preceding January

```
In [18]: # Group off-peak prices by companies and month
monthly_price_by_id = price_data.groupby(['id', 'price_date']).agg({'price_off_peak_var': 'mean', 'price_off_peak_fix': 'mean'})

# Get january and december prices
jan_prices = monthly_price_by_id.groupby('id').first().reset_index()
dec_prices = monthly_price_by_id.groupby('id').last().reset_index()

# Calculate the difference
diff = pd.merge(dec_prices.rename(columns={'price_off_peak_var': 'dec_1', 'price_off_peak_fix': 'dec_2'}), jan_prices.rename(columns={'price_off_peak_var': 'jan_1', 'price_off_peak_fix': 'jan_2'}), on='id')
diff['offpeak_diff_dec_january_energy'] = diff['dec_1'] - diff['jan_1']
diff['offpeak_diff_dec_january_power'] = diff['dec_2'] - diff['jan_2']
diff = diff[['id', 'offpeak_diff_dec_january_energy', 'offpeak_diff_dec_january_power']]
diff.head()
```

```
Out[18]:
```

	id	offpeak_diff_dec_january_energy	offpeak_diff_dec_january_power
0	0002203ffb812588b632b9e628cc38d	-0.006192	0.162916
1	0004351ebdd665e6ee664792efc4fd13	-0.004104	0.177779
2	0010bcc39e42b3c2131ed2ce55246e3c	0.050443	1.500000
3	0010ee3855fdea87602a5b7aba8e42de	-0.010018	0.162916
4	00114d74e963e47177db89bc70108537	-0.003994	-0.000001

```
In [19]: df = pd.merge(client_data, diff, on='id')
df.head()
```

```
Out[19]:
```

	id	channel_sales	cons_12m	cons_gas_12m	cons_last_month	date_activ
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwksosbicdxkicaua	0	54946	0	2013-06-15
1	d29c2c54acc38ff3c0614d0a653813dd	MISSING	4660	0	0	2009-08-21
2	764c75f661154dac3a6c254cd082ea7d	foosdfpfkusacimwksosbicdxkicaua	544	0	0	2010-04-16
3	bba03439a292a1e166f80264c16191cb	lmkebamcaaclubfxadlmueccxoimlema	1584	0	0	2010-03-30
4	149d57cf92fc41cf94415803a877cb4b	MISSING	4425	0	526	2010-01-13

5 rows × 28 columns

Average price changes across periods

```
In [20]: # Aggregate average prices per period by company
mean_prices = price_data.groupby(['id']).agg({
    'price_off_peak_var': 'mean',
    'price_peak_var': 'mean',
    'price_mid_peak_var': 'mean',
    'price_off_peak_fix': 'mean',
    'price_peak_fix': 'mean',
    'price_mid_peak_fix': 'mean'
}).reset_index()
```

```
In [21]: # Calculate the mean difference between consecutive periods
mean_prices['off_peak_peak_var_mean_diff'] = mean_prices['price_off_peak_var'] - mean_prices['price_peak_var']
mean_prices['peak_mid_peak_var_mean_diff'] = mean_prices['price_peak_var'] - mean_prices['price_mid_peak_var']
mean_prices['off_peak_mid_peak_var_mean_diff'] = mean_prices['price_off_peak_var'] - mean_prices['price_mid_peak_var']
mean_prices['off_peak_peak_fix_mean_diff'] = mean_prices['price_off_peak_fix'] - mean_prices['price_peak_fix']
mean_prices['peak_mid_peak_fix_mean_diff'] = mean_prices['price_peak_fix'] - mean_prices['price_mid_peak_fix']
mean_prices['off_peak_mid_peak_fix_mean_diff'] = mean_prices['price_off_peak_fix'] - mean_prices['price_mid_peak_fix']
```

```
In [22]: columns = ['id', 'off_peak_peak_var_mean_diff', 'peak_mid_peak_var_mean_diff', 'off_peak_mid_peak_var_mean_diff']
df = pd.merge(df, mean_prices[columns], on='id')
df.head()
```

	id	channel_sales	cons_12m	cons_gas_12m	cons_last_month	date_activ
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwksosbicdxkicaua	0	54946	0	2013-06-15
1	d29c2c54acc38ff3c0614d0a653813dd	MISSING	4660	0	0	2009-08-21
2	764c75f661154dac3a6c254cd082ea7d	foosdfpfkusacimwksosbicdxkicaua	544	0	0	2010-04-16
3	bba03439a292a1e166f80264c16191cb	lmkebamcaaclubfxadlmueccxoimlema	1584	0	0	2010-03-30
4	149d57cf92fc41cf94415803a877cb4b	MISSING	4425	0	526	2010-01-13

```
# Aggregate average prices per period by company
mean_prices_by_month = price_data.groupby(['id', 'price_date']).agg({'price_off_peak_var': 'mean', 'price_peak_
```

```
# Calculate the mean difference between consecutive periods
mean_prices_by_month['off_peak_peak_var_mean_diff'] = mean_prices_by_month['price_off_peak_var'] - mean_prices_by_month['price_peak_var']
mean_prices_by_month['peak_mid_peak_var_mean_diff'] = mean_prices_by_month['price_peak_var'] - mean_prices_by_month['price_mid_peak_var']
mean_prices_by_month['off_peak_mid_peak_var_mean_diff'] = mean_prices_by_month['price_off_peak_var'] - mean_prices_by_month['price_mid_peak_var']
mean_prices_by_month['off_peak_peak_fix_mean_diff'] = mean_prices_by_month['price_off_peak_fix'] - mean_prices_by_month['price_peak_fix']
mean_prices_by_month['peak_mid_peak_fix_mean_diff'] = mean_prices_by_month['price_peak_fix'] - mean_prices_by_month['price_mid_peak_fix']
mean_prices_by_month['off_peak_mid_peak_fix_mean_diff'] = mean_prices_by_month['price_off_peak_fix'] - mean_prices_by_month['price_mid_peak_fix']
```

```
# Calculate the maximum monthly difference across time periods
max_diff_across_periods_months = mean_prices_by_month.groupby(['id']).agg({'off_peak_peak_var_mean_diff': 'max'})
```

```
columns = ['id', 'off_peak_peak_var_max_monthly_diff', 'peak_mid_peak_var_max_monthly_diff', 'off_peak_mid_peak_va
df = pd.merge(df, max_diff_across_periods_months[columns], on='id')
df.head()
```

	id	channel_sales	cons_12m	cons_gas_12m	cons_last_month	date_activ
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwksosbicdxkicaua	0	54946	0	2013-06-15
1	d29c2c54acc38ff3c0614d0a653813dd	MISSING	4660	0	0	2009-08-21
2	764c75f661154dac3a6c254cd082ea7d	foosdfpfkusacimwksosbicdxkicaua	544	0	0	2010-04-16
3	bba03439a292a1e166f80264c16191cb	lmkebamcaaclubfxadlmueccxoimlema	1584	0	0	2010-03-30
4	149d57cf92fc41cf94415803a877cb4b	MISSING	4425	0	526	2010-01-13

```
df['tenure'] = df.date_end.dt.year - df.date_activ.dt.year
```

```
df.groupby(['tenure']).agg({'churn': 'mean'}).sort_values(by='churn', ascending=False)
```


Out[28]:

	churn
tenure	
3	0.144612
4	0.126383
5	0.099897
13	0.093023
12	0.085106
6	0.075593
7	0.075025
8	0.058065
11	0.052356
9	0.037037
10	0.027778
2	0.000000

Transforming dates into months

- months_activ = Number of months active until reference date (Jan 2016)
- months_to_end = Number of months of the contract left until reference date (Jan 2016)
- months_modif_prod = Number of months since last modification until reference date (Jan 2016)
- months_renewal = Number of months since last renewal until reference date (Jan 2016)

```
In [29]: def convert_months(reference_date, df, column):
         months = (reference_date.year - df[column].dt.year)*12 + (reference_date.month - df[column].dt.month)
         return months
```

```
In [30]: # Create reference date
reference_date = pd.to_datetime('2016-01-01')

# Create columns
df['months_activ'] = convert_months(reference_date, df, 'date_activ')
df['months_to_end'] = -convert_months(reference_date, df, 'date_end')
df['months_modif_prod'] = convert_months(reference_date, df, 'date_modif_prod')
df['months_renewal'] = convert_months(reference_date, df, 'date_renewal')
```

```
In [31]: remove = ['date_activ', 'date_end', 'date_modif_prod', 'date_renewal']
df = df.drop(columns=remove)
df.head()
```

Out[31]:

	id	channel_sales	cons_12m	cons_gas_12m	cons_last_month	forecast_co
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsosbicdxkicaua	0	54946	0	
1	d29c2c54acc38ff3c0614d0a653813dd	MISSING	4660	0	0	
2	764c75f661154dac3a6c254cd082ea7d	foosdfpfkusacimwkcsosbicdxkicaua	544	0	0	
3	bba03439a292a1e166f80264c16191cb	lmkebamcaaclubfxadlmueccxoimlema	1584	0	0	
4	149d57cf92fc41cf94415803a877cb4b	MISSING	4425	0	526	

5 rows × 41 columns

--	--	--

```
In [32]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14606 entries, 0 to 14605
Data columns (total 41 columns):
```

#	Column	Non-Null Count	Dtype
0	id	14606 non-null	object
1	channel_sales	14606 non-null	object
2	cons_12m	14606 non-null	int64
3	cons_gas_12m	14606 non-null	int64
4	cons_last_month	14606 non-null	int64
5	forecast_cons_12m	14606 non-null	float64
6	forecast_cons_year	14606 non-null	int64
7	forecast_discount_energy	14606 non-null	float64
8	forecast_meter_rent_12m	14606 non-null	float64
9	forecast_price_energy_off_peak	14606 non-null	float64
10	forecast_price_energy_peak	14606 non-null	float64
11	forecast_price_pow_off_peak	14606 non-null	float64
12	has_gas	14606 non-null	object
13	imp_cons	14606 non-null	float64
14	margin_gross_pow_ele	14606 non-null	float64
15	margin_net_pow_ele	14606 non-null	float64
16	nb_prod_act	14606 non-null	int64
17	net_margin	14606 non-null	float64
18	num_years_antig	14606 non-null	int64
19	origin_up	14606 non-null	object
20	pow_max	14606 non-null	float64
21	churn	14606 non-null	int64
22	offpeak_diff_dec_january_energy	14606 non-null	float64
23	offpeak_diff_dec_january_power	14606 non-null	float64
24	off_peak_peak_var_mean_diff	14606 non-null	float64
25	peak_mid_peak_var_mean_diff	14606 non-null	float64
26	off_peak_mid_peak_var_mean_diff	14606 non-null	float64
27	off_peak_peak_fix_mean_diff	14606 non-null	float64
28	peak_mid_peak_fix_mean_diff	14606 non-null	float64
29	off_peak_mid_peak_fix_mean_diff	14606 non-null	float64
30	off_peak_peak_var_max_monthly_diff	14606 non-null	float64
31	peak_mid_peak_var_max_monthly_diff	14606 non-null	float64
32	off_peak_mid_peak_var_max_monthly_diff	14606 non-null	float64
33	off_peak_peak_fix_max_monthly_diff	14606 non-null	float64
34	peak_mid_peak_fix_max_monthly_diff	14606 non-null	float64
35	off_peak_mid_peak_fix_max_monthly_diff	14606 non-null	float64
36	tenure	14606 non-null	int32
37	months_activ	14606 non-null	int32
38	months_to_end	14606 non-null	int32
39	months_modif_prod	14606 non-null	int32
40	months_renewal	14606 non-null	int32

dtypes: float64(25), int32(5), int64(7), object(4)
memory usage: 4.3+ MB

```
In [33]: df = pd.get_dummies(df, columns=["origin_up"], prefix="origin_up", dtype = int)
df
```

```
Out[33]:
```

	id	channel_sales	cons_12m	cons_gas_12m	cons_last_month	forecas
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsoibcdxkicaua	0	54946	0	
1	d29c2c54acc38ff3c0614d0a653813dd	MISSING	4660	0	0	
2	764c75f661154dac3a6c254cd082ea7d	foosdfpfkusacimwkcsoibcdxkicaua	544	0	0	
3	bba03439a292a1e166f80264c16191cb	lmkebamcaaclubfxadlmueccxoimlema	1584	0	0	
4	149d57cf92fc41cf94415803a877cb4b	MISSING	4425	0	526	
...
14601	18463073fb097fc0ac5d3e040f356987	foosdfpfkusacimwkcsoibcdxkicaua	32270	47940	0	
14602	d0a6f71671571ed83b2645d23af6de00	foosdfpfkusacimwkcsoibcdxkicaua	7223	0	181	
14603	10e6828ddd62cbcf687cb74928c4c2d2	foosdfpfkusacimwkcsoibcdxkicaua	1844	0	179	
14604	1cf20fd6206d7678d5bcafd28c53b4db	foosdfpfkusacimwkcsoibcdxkicaua	131	0	0	
14605	563dde550fd624d7352f3de77c0cdfcd	MISSING	8730	0	0	

14606 rows × 46 columns




```
In [34]: df = pd.get_dummies(df, columns=["channel_sales"], prefix="channel_sales", dtype = int)
df
```

Out [34]:

	id	cons_12m	cons_gas_12m	cons_last_month	forecast_cons_12m	forecast_cons_year	fc
0	24011ae4ebbe3035111d65fa7c15bc57	0	54946	0	0.00	0	
1	d29c2c54acc38ff3c0614d0a653813dd	4660	0	0	189.95	0	
2	764c75f661154dac3a6c254cd082ea7d	544	0	0	47.96	0	
3	bba03439a292a1e166f80264c16191cb	1584	0	0	240.04	0	
4	149d57cf92fc41cf94415803a877cb4b	4425	0	526	445.75	526	
...
14601	18463073fb097fc0ac5d3e040f356987	32270	47940	0	4648.01	0	
14602	d0a6f71671571ed83b2645d23af6de00	7223	0	181	631.69	181	
14603	10e6828ddd62cbcf687cb74928c4c2d2	1844	0	179	190.39	179	
14604	1cf20fd6206d7678d5bcafd28c53b4db	131	0	0	19.34	0	
14605	563dde550fd624d7352f3de77c0cdfcd	8730	0	0	762.41	0	

14606 rows × 53 columns



In [35]:

```
df.has_gas.replace({'f': 0, 't': 1}, inplace=True)
df
```

Out [35]:

	id	cons_12m	cons_gas_12m	cons_last_month	forecast_cons_12m	forecast_cons_year	fc
0	24011ae4ebbe3035111d65fa7c15bc57	0	54946	0	0.00	0	
1	d29c2c54acc38ff3c0614d0a653813dd	4660	0	0	189.95	0	
2	764c75f661154dac3a6c254cd082ea7d	544	0	0	47.96	0	
3	bba03439a292a1e166f80264c16191cb	1584	0	0	240.04	0	
4	149d57cf92fc41cf94415803a877cb4b	4425	0	526	445.75	526	
...
14601	18463073fb097fc0ac5d3e040f356987	32270	47940	0	4648.01	0	
14602	d0a6f71671571ed83b2645d23af6de00	7223	0	181	631.69	181	
14603	10e6828ddd62cbcf687cb74928c4c2d2	1844	0	179	190.39	179	
14604	1cf20fd6206d7678d5bcafd28c53b4db	131	0	0	19.34	0	
14605	563dde550fd624d7352f3de77c0cdfcd	8730	0	0	762.41	0	

14606 rows × 53 columns



In [36]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14606 entries, 0 to 14605
Data columns (total 53 columns):
```

#	Column	Non-Null Count	Dtype
0	id	14606 non-null	object
1	cons_12m	14606 non-null	int64
2	cons_gas_12m	14606 non-null	int64
3	cons_last_month	14606 non-null	int64
4	forecast_cons_12m	14606 non-null	float64
5	forecast_cons_year	14606 non-null	int64
6	forecast_discount_energy	14606 non-null	float64
7	forecast_meter_rent_12m	14606 non-null	float64
8	forecast_price_energy_off_peak	14606 non-null	float64
9	forecast_price_energy_peak	14606 non-null	float64
10	forecast_price_pow_off_peak	14606 non-null	float64
11	has_gas	14606 non-null	int64
12	imp_cons	14606 non-null	float64
13	margin_gross_pow_ele	14606 non-null	float64
14	margin_net_pow_ele	14606 non-null	float64
15	nb_prod_act	14606 non-null	int64
16	net_margin	14606 non-null	float64
17	num_years_antig	14606 non-null	int64
18	pow_max	14606 non-null	float64
19	churn	14606 non-null	int64
20	offpeak_diff_dec_january_energy	14606 non-null	float64
21	offpeak_diff_dec_january_power	14606 non-null	float64
22	off_peak_peak_var_mean_diff	14606 non-null	float64
23	peak_mid_peak_var_mean_diff	14606 non-null	float64
24	off_peak_mid_peak_var_mean_diff	14606 non-null	float64
25	off_peak_peak_fix_mean_diff	14606 non-null	float64
26	peak_mid_peak_fix_mean_diff	14606 non-null	float64
27	off_peak_mid_peak_fix_mean_diff	14606 non-null	float64
28	off_peak_peak_var_max_monthly_diff	14606 non-null	float64
29	peak_mid_peak_var_max_monthly_diff	14606 non-null	float64
30	off_peak_mid_peak_var_max_monthly_diff	14606 non-null	float64
31	off_peak_peak_fix_max_monthly_diff	14606 non-null	float64
32	peak_mid_peak_fix_max_monthly_diff	14606 non-null	float64
33	off_peak_mid_peak_fix_max_monthly_diff	14606 non-null	float64
34	tenure	14606 non-null	int32
35	months_activ	14606 non-null	int32
36	months_to_end	14606 non-null	int32
37	months_modif_prod	14606 non-null	int32
38	months_renewal	14606 non-null	int32
39	origin_up_MISSING	14606 non-null	int32
40	origin_up_ewxeelcelemmiuafmddpobolxfuxioce	14606 non-null	int32
41	origin_up_kamkxfxxuwbdsllkwifmmcsiusiuosws	14606 non-null	int32
42	origin_up_ldkssxwpmemidmecebumciepifcamkci	14606 non-null	int32
43	origin_up_lxidpiddsbxsbsoboudacockeimpuepw	14606 non-null	int32
44	origin_up_usapbecpcfoloekilkwdsiboslwaxobdp	14606 non-null	int32
45	channel_sales_MISSING	14606 non-null	int32
46	channel_sales_epumfxlbckeskwexbiuasklxalciuu	14606 non-null	int32
47	channel_sales_ewpakwlliwisiwduibdlfmalxowmwpci	14606 non-null	int32
48	channel_sales_fixdbufsefwooaasfcxdxadsiekoceaa	14606 non-null	int32
49	channel_sales_foosdfpfkusacimwkcsosbicdxkicaua	14606 non-null	int32
50	channel_sales_lmkebamcaaclubfxadlmueccxoimlema	14606 non-null	int32
51	channel_sales_sddiedcslfslkckwlfdkpoeaailfpeds	14606 non-null	int32
52	channel_sales_usilxuppasemubllpokaafesmlibmsdf	14606 non-null	int32

```
dtypes: float64(25), int32(19), int64(8), object(1)
memory usage: 4.8+ MB
```

Modeling the Data

```
In [37]: X = df.drop(['id', 'churn'], axis=1)
X.shape
```

```
Out[37]: (14606, 51)
```

```
In [38]: y = df.churn
y.shape
```

```
Out[38]: (14606,)
```

Splitting dataset into Training and Testing

```
In [39]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)
```

```
In [40]: print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(11684, 51)
(2922, 51)
(11684,)
(2922,)
```

Scaling the Data


```
In [41]: scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```


Building the ANN Model


```
In [42]: model = Sequential([
    Flatten(),
    Dense(45, activation='relu'),
    Dropout(0.2),
    Dense(35, activation='relu'),
    Dense(25, activation='relu'),
    Dense(10, activation='relu'),
    Dense(1, activation='sigmoid'),
])
```


```
In [43]: model.compile(
    optimizer = 'adam',
    loss = 'binary_crossentropy',
    metrics = ['accuracy']
)
```


```
In [44]: history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=20, batch_size=32)
```


Epoch 1/20
366/366  **3s** 3ms/step - accuracy: 0.8793 - loss: 0.3973 - val_accuracy: 0.9028 - val_loss: 0.3105


Epoch 2/20
366/366  **1s** 2ms/step - accuracy: 0.9050 - loss: 0.3088 - val_accuracy: 0.9028 - val_loss: 0.3128


Epoch 3/20
366/366  **1s** 2ms/step - accuracy: 0.9000 - loss: 0.3177 - val_accuracy: 0.9028 - val_loss: 0.3093


Epoch 4/20
366/366  **1s** 2ms/step - accuracy: 0.9082 - loss: 0.2989 - val_accuracy: 0.9028 - val_loss: 0.3085


Epoch 5/20
366/366  **1s** 2ms/step - accuracy: 0.9016 - loss: 0.3137 - val_accuracy: 0.9028 - val_loss: 0.3086


Epoch 6/20
366/366  **1s** 2ms/step - accuracy: 0.9042 - loss: 0.3033 - val_accuracy: 0.9028 - val_loss: 0.3073


Epoch 7/20
366/366  **1s** 2ms/step - accuracy: 0.9068 - loss: 0.2968 - val_accuracy: 0.9028 - val_loss: 0.3147


Epoch 8/20
366/366  **1s** 2ms/step - accuracy: 0.9061 - loss: 0.3008 - val_accuracy: 0.9028 - val_loss: 0.3065


Epoch 9/20
366/366  **1s** 2ms/step - accuracy: 0.9064 - loss: 0.2977 - val_accuracy: 0.9028 - val_loss: 0.3069


Epoch 10/20
366/366  **1s** 2ms/step - accuracy: 0.9036 - loss: 0.3018 - val_accuracy: 0.9028 - val_loss: 0.3070


Epoch 11/20
366/366  **1s** 3ms/step - accuracy: 0.9088 - loss: 0.2905 - val_accuracy: 0.9028 - val_loss: 0.3099


Epoch 12/20
366/366  **1s** 3ms/step - accuracy: 0.9052 - loss: 0.2998 - val_accuracy: 0.9028 - val_loss: 0.3074


Epoch 13/20
366/366  **1s** 2ms/step - accuracy: 0.9066 - loss: 0.2944 - val_accuracy: 0.9028 - val_loss: 0.3088


Epoch 14/20
366/366  **1s** 3ms/step - accuracy: 0.9008 - loss: 0.3079 - val_accuracy: 0.9028 - val_loss: 0.3088


Epoch 15/20
366/366  **1s** 2ms/step - accuracy: 0.9040 - loss: 0.2986 - val_accuracy: 0.9028 - val_loss: 0.3046

Epoch 16/20
366/366  **1s** 3ms/step - accuracy: 0.9014 - loss: 0.3051 - val_accuracy: 0.9028 - val_loss: 0.3072

Epoch 17/20
366/366  **1s** 2ms/step - accuracy: 0.9064 - loss: 0.2894 - val_accuracy: 0.9028 - val_loss: 0.3074

Epoch 18/20
366/366  **1s** 3ms/step - accuracy: 0.8975 - loss: 0.3100 - val_accuracy: 0.9028 - val_loss: 0.3069

Epoch 19/20
366/366  **1s** 2ms/step - accuracy: 0.8977 - loss: 0.3110 - val_accuracy: 0.9028 - val_loss: 0.3078

Epoch 20/20
366/366  **1s** 2ms/step - accuracy: 0.9064 - loss: 0.2891 - val_accuracy: 0.9028 - val_loss: 0.3070


```
In [45]: y_pred_values = model.predict(X_test)
```

92/92  **0s** 2ms/step

```
In [46]: y_pred = []
for value in y_pred_values:
    if value > 0.5:
        y_pred.append(1)
    else:
        y_pred.append(0)
```

Model Evaluation

```
In [47]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='best')
plt.show()
```



```
In [48]: model.evaluate(X_test, y_test)
```

92/92 ————— 0s 1ms/step - accuracy: 0.9005 - loss: 0.3118

```
Out[48]: [0.30697688460350037, 0.902806282043457]
```

```
In [49]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.90	1.00	0.95	2638
1	0.00	0.00	0.00	284
accuracy			0.90	2922
macro avg	0.45	0.50	0.47	2922
weighted avg	0.82	0.90	0.86	2922

```
In [50]: sns.heatmap(confusion_matrix(y_test, y_pred), annot=True)
plt.title(f"Accuracy for ANN: {round(accuracy_score(y_test, y_pred),2)}")
plt.xlabel("Predicted")
plt.ylabel("Truth")
plt.show()
```

